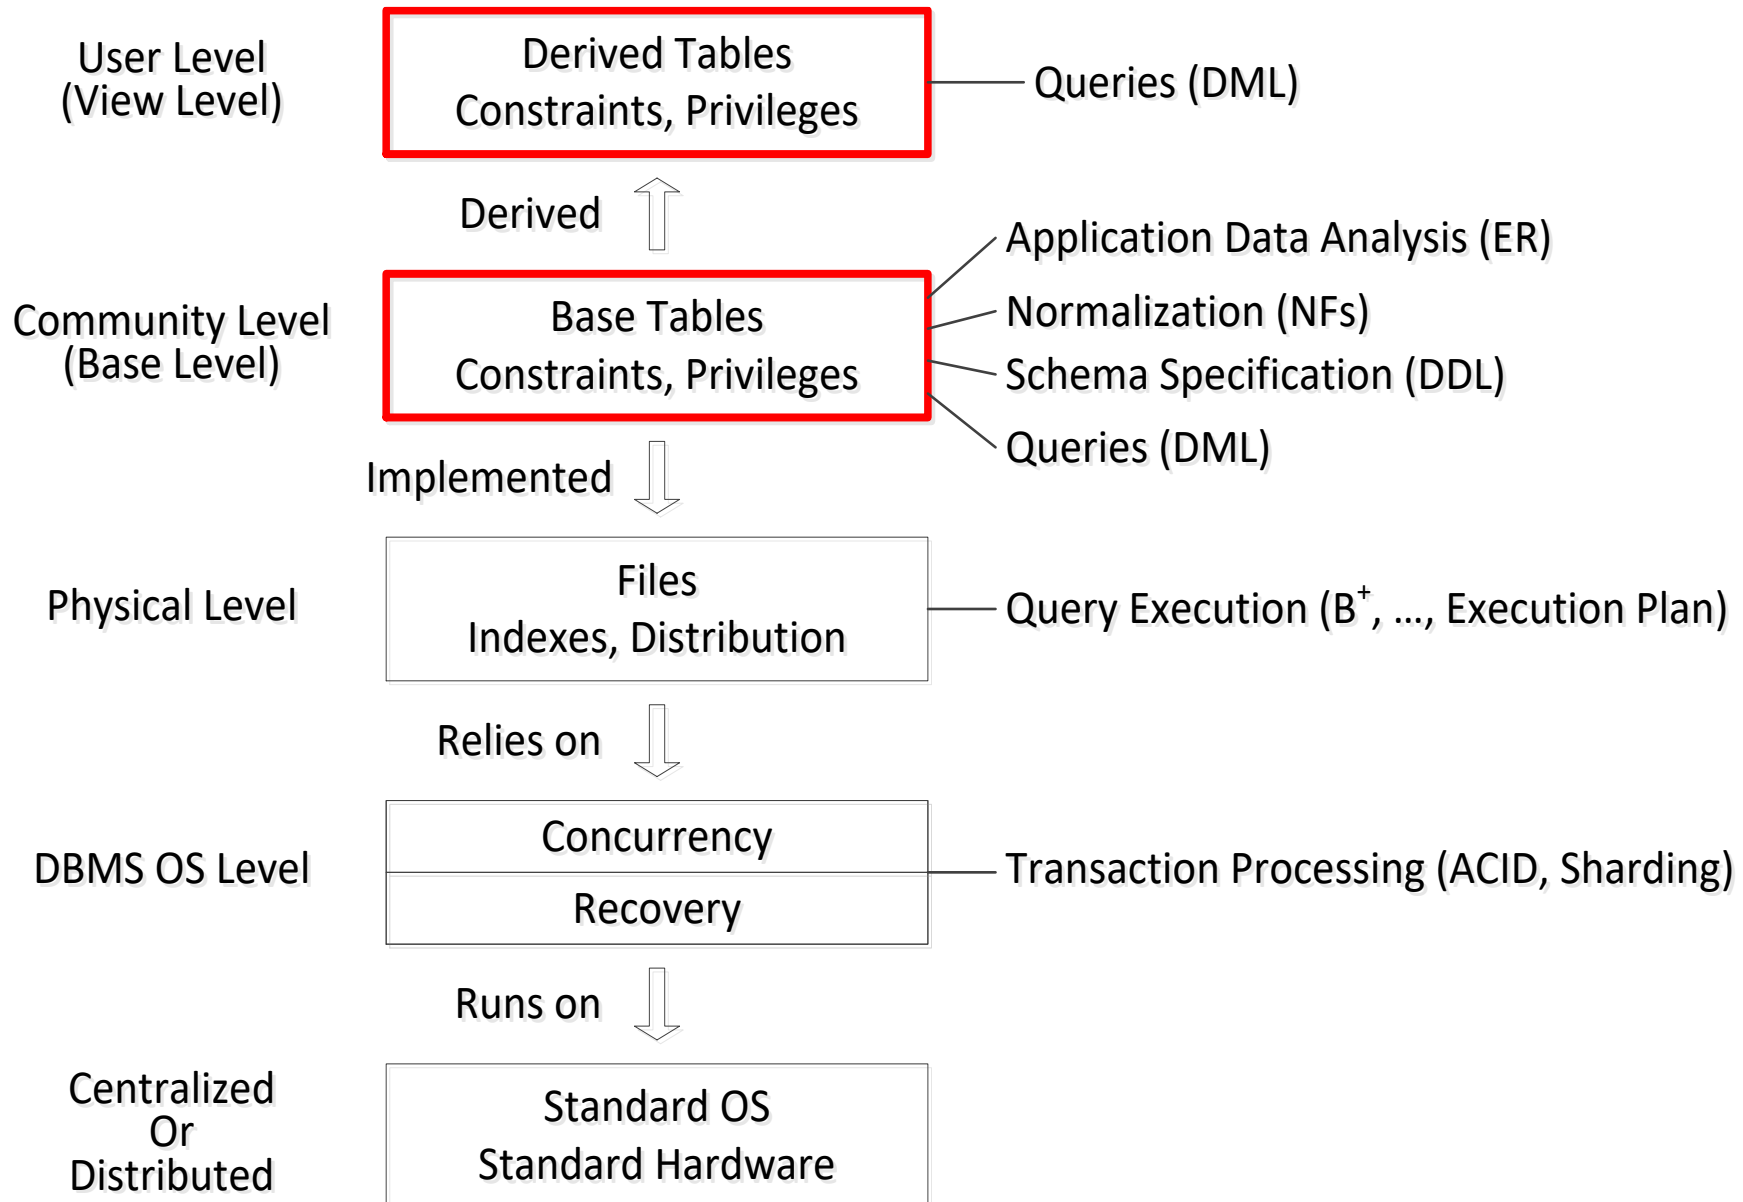


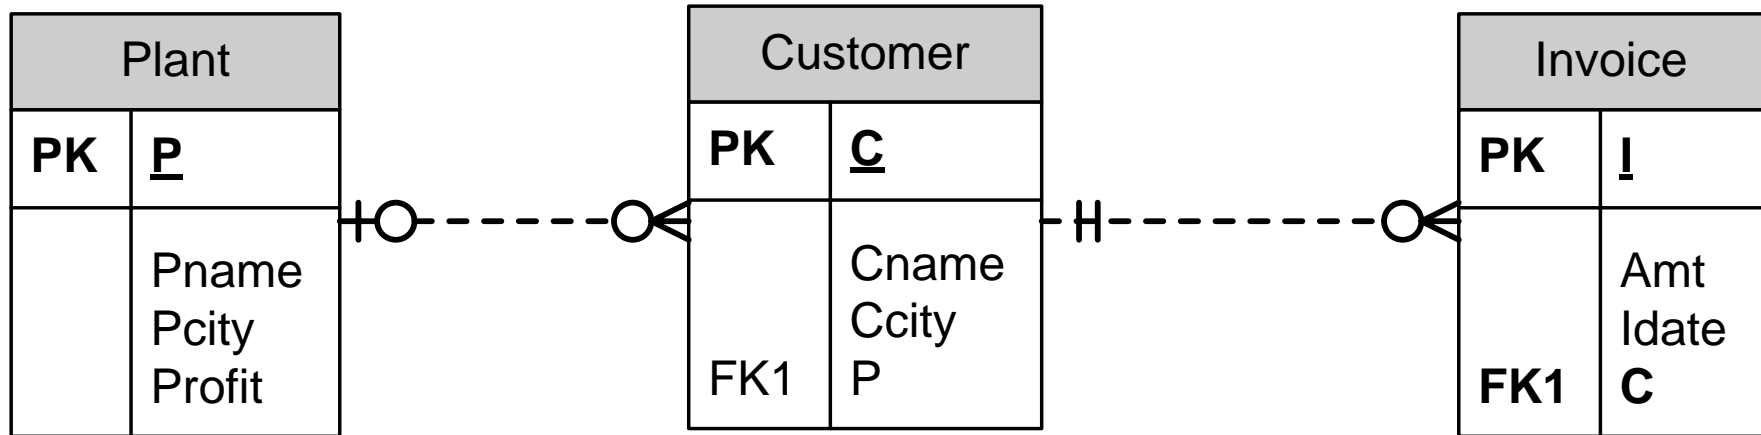
***Unit 6***  
***SQL: Data Definition Language***  
***And Data Control Language***  
***For Relational Databases***

# DDL and DCL in Context



# ***Introduction***

# *The Tables To Be Defined And Some More*



- This is the database we will define
- We do not pay attention to domains of attributes as there is not much interesting in that
- |-0 means zero or one (customer has zero or one plant). |-| means exactly one (invoice goes to one customer).

# ***Defining A Relational Database***

- We will focus only on some of the basic capabilities for defining a relational database
- The standard is very extensive and provides for a rich repertoire of useful capabilities
- We touch on the basics.
- But enough for defining reasonable-complexity databases

***CREATE TABLE***  
***(Specifying A Relation Schema)***

## ***Basic Definition***

- CREATE TABLE Plant (  
    P CHAR(10),  
    Pname CHAR VARYING(10),  
    Pcity CHAR VARYING(10),  
    Profit NUMBER  
);
  
- This is a minimal definition
  - Name of the table
  - Names of the columns
  - Domains of the columns

# ***Basic Definition***

- ```
CREATE TABLE Customer (  
  C CHAR(10),  
  Cname CHAR VARYING(10),  
  Ccity CHAR VARYING(10),  
  P CHAR(10)  
);
```

- This is a minimal definition
  - Name of the table
  - Names of the columns
  - Domains of the columns



# ***Basic Definition***

- `CREATE TABLE Invoice (  
    I CHAR(10),  
    Amt NUMBER,  
    Idate DATE,  
    C CHAR(10)  
);`

- This is a minimal definition
  - Name of the table
  - Names of the columns
  - Domains of the columns

# ***Permitted Data Types (Data Domains)***

- SQL standard specifies permitted data types, which can be roughly grouped into several families
  - Integers (small or long)
  - Real numbers (standard or double length and with various precisions)
  - Character strings (fixed or variable length)
  - Bit strings (fixed or variable length)
  - Dates and times (various specifications with various time “granularity”)
- Systems have different implementations and modifications of the standard

# ***Notation***

- In some of the slides, new concepts will be introduced
- The SQL specifications will be in red color and bold to draw attention to them

## *Minimum Specification For Plant*

- **CREATE TABLE Plant (**  
    **P CHAR(10) NOT NULL,**  
    **Pname CHAR VARYING(10),**  
    **Pcity CHAR VARYING(10),**  
    **Profit NUMBER,**  
    **CONSTRAINT C\_20 PRIMARY KEY (P),**  
    **CONSTRAINT C\_30 UNIQUE (Pcity, Profit),**  
    **CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),**  
    **CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR**  
    **(Profit > 1000) )**  
    **);**
- This is a minimal definition
  - Name of the table
  - Names of the columns
  - Domains of the columns

## ***Not Null***

- CREATE TABLE Plant (  
P CHAR(10) **NOT NULL**,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 PRIMARY KEY (P),  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR  
(Profit > 1000) )  
);
- Specifies that the values in these columns (could be more than one such column) must not be NULL

# Constraints

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
**CONSTRAINT C\_20** PRIMARY KEY (P),  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR  
(Profit > 1000) )  
);
- Some constraint on the tables
  - Constraint name, here C\_20, is not required, but it is a very good idea to give unique names to a constraint, so it can be later DROPPed or ALTERed by referring to it by its name
  - Constraint name should reflect something about the constraint, but to save space I used short names

# ***Primary Key***

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 **PRIMARY KEY (P)**,  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR  
(Profit > 1000) )  
);
- The column P is the primary key (only one possible)
  - This requires that it must not be NULL (this is not necessary to state in some systems, as the primary key condition automatically forces it by SQL standard)
- Primary key could be several columns, e.g., PRIMARY KEY(Pcity, Profit); but not in our example

# Unique

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 PRIMARY KEY (P),  
CONSTRAINT C\_30 **UNIQUE (Pcity, Profit)**,  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR  
(Profit > 1000) )  
);
- The “subtuple” Pcity,Pname is a candidate key
  - There is no requirement, in general, about any of its column being not NULL
  - But for every value of the pair there is at most one tuple in Plant
  - To reiterate: all the columns of the primary key must not be NULL



## ***Check (And Unknown)***

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 PRIMARY KEY (P),  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 **CHECK ( Pcity <> Pname )**,  
CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR  
(Profit > 1000) )  
);
- Every tuple must satisfy this condition
- The condition is **satisfied**, when it is either
  - **TRUE, or**
  - **UNKNOWN** (so if Pcity is Null, this condition is satisfied)
- **Recall in SQL SELECT queries: UNKNOWN implies not satisfied, effectively FALSE; here effectively TRUE**

# Check

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 PRIMARY KEY (P),  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( (**Pcity <> 'Chicago'**) OR  
**(Profit > 1000)** )  
);
- This is: (Pcity = 'Chicago') → (Profit > 1000)  
By standard rules of Boolean operators (propositional calculus)

## ***Check (And Unknown)***

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
CONSTRAINT C\_20 PRIMARY KEY (P),  
CONSTRAINT C\_30 UNIQUE (Pcity, Profit),  
CONSTRAINT C\_40 CHECK ( Pcity <> Pname ),  
CONSTRAINT C\_50 CHECK ( **(Pcity <> 'Chicago') OR**  
**(Profit > 1000) )**  
);
- Returning to semantics of UNKNOWN and OR, this constraint has to evaluate to TRUE or UNKNOWN to be satisfied, so we need  
( Pcity is not Chicago or IS NULL) or ( Profit is greater than 1000 or IS NULL)
- So for Chicago the profit is greater than 1000 or IS NULL

## ***Defaults***

- CREATE TABLE Customer (  
C CHAR(10) NOT NULL,  
Cname CHAR VARYING(10) **DEFAULT NULL**,  
Ccity CHAR VARYING(10),  
P CHAR(10) **DEFAULT 'Main'**,  
CONSTRAINT C\_60 PRIMARY KEY (C),  
CONSTRAINT C\_70 FOREIGN KEY (P) REFERENCES  
Plant(P) ON DELETE SET NULL  
);
- It is possible to specify defaults
  - E.g., when a tuple is inserted and only C and Ccity are specified, the system knows to specify NULL for Cname and Main for P

# Foreign Key

- CREATE TABLE Customer (  
C CHAR(10) NOT NULL,  
Cname CHAR VARYING(10) DEFAULT NULL,  
Ccity CHAR VARYING(10),  
P CHAR(10) DEFAULT 'Main',  
CONSTRAINT C\_60 PRIMARY KEY (C),  
CONSTRAINT C\_70 **FOREIGN KEY (P) REFERENCES**  
**Plant(P)** ON DELETE SET NULL  
);
- P in Customer has to reference the primary key of Plant (in this example) or a UNIQUE (not in this example)
- This means that one of two conditions is satisfied
  - P has a non NULL value and this value of P appears in Plant
  - P is NULL

Of course, if P were specified as NOT NULL, this could not be the case

## ***On Delete Set Null***

- CREATE TABLE Customer (  
C CHAR(10) NOT NULL,  
Cname CHAR VARYING(10) DEFAULT NULL,  
Ccity CHAR VARYING(10),  
P CHAR(10) DEFAULT 'Main',  
CONSTRAINT C\_60 PRIMARY KEY (C),  
CONSTRAINT C\_70 FOREIGN KEY (P) REFERENCES  
Plant(P) **ON DELETE SET NULL**  
);
- P in Customer has to reference the primary key of Plant or a UNIQUE
- But note, that P in Customer is not required to be NOT NULL
- We have a specification that if P is listed in some tuple of Customer and that tuple with that P is deleted from Plant then that value of P in Customer is replaced by NULL

## ***Not Null***

- CREATE TABLE Invoice (  
I CHAR(10) NOT NULL,  
Amt NUMBER,  
Idate DATE,  
C CHAR(10) **NOT NULL**,  
CONSTRAINT C\_80 PRIMARY KEY (I),  
CONSTRAINT C\_90 FOREIGN KEY (C) REFERENCES  
Customer(C) ON DELETE CASCADE  
);
- NOT NULL can be specified also for columns not in the primary key

## ***On Delete Cascade***

- CREATE TABLE Invoice (  
  I CHAR(10) NOT NULL,  
  Amt NUMBER,  
  Idate DATE,  
  C CHAR(10) NOT NULL,  
  CONSTRAINT C\_80 PRIMARY KEY (I),  
  CONSTRAINT C\_90 FOREIGN KEY (C) REFERENCES  
  Customer(C) **ON DELETE CASCADE**  
);
- We have a specification that if C listed in some tuple of Invoice is deleted from Customer (that is the tuple with this value of primary key is deleted), all the tuples with this value of C in Invoice must be deleted



# ***Maintenance of Referential Integrity***

- In order to maintain referential integrity constraints, the system will reject any operation that will violate it.
  - There are subtle interactions if NULLs are present; we will not discuss them here
- CREATE TABLE Invoice (  
I CHAR(10) NOT NULL,  
Amt NUMBER,  
Idate DATE,  
C CHAR(10) NOT NULL,  
CONSTRAINT C\_80 PRIMARY KEY (I),  
CONSTRAINT C\_90 **FOREIGN KEY (C) REFERENCES**  
**Customer(C) ON ...**  
);

# ***Maintenance of Referential Integrity On Update***

- This constraint “will act” when:
  - An **INSERT** or an **UPDATE** on Invoice is attempted that would produce there a value of C that does not exist in Customer.
  - A **DELETE** or an **UPDATE** on Customer is attempted that will leave tuples in Invoice in which the value of C does not appear in any tuple of Customer.
- The default is **NO ACTION**, that is the above will not be permitted
- We will briefly discuss other options in case of UPDATES of Customer and skip what happens in other cases
  - **CASCADE**: the new value of the primary key (or UNIQUE if the foreign key refers to UNIQUE attribute or attributes) is copied to the foreign key
  - **SET NULL**: the new value of the foreign key (or UNIQUE) is NULL (if NULL is allowed)
  - **SET DEFAULT**: the new value of the foreign key (or UNIQUE) is a specified default value (which of course has to appear in Customer)

# ***Starting With A Basic Definition***

- It is generally a good idea to start with a basic definition and augment it with constraints later
- We see how this is done

## ***Basic Definition***

- CREATE TABLE Plant (  
    P CHAR(10) NOT NULL,  
    Pname CHAR VARYING(10),  
    Pcity CHAR VARYING(10),  
    Profit NUMBER  
);

## ***Basic Definition***

- CREATE TABLE Customer (  
    C CHAR(10) NOT NULL,  
    Cname CHAR VARYING(10) DEFAULT (NULL),  
    Ccity CHAR VARYING(10),  
    P CHAR(10) DEFAULT ('Main')  
);

## ***Basic Definition***

- CREATE TABLE Invoice (  
  I CHAR(10) NOT NULL,  
  Amt NUMBER,  
  Idate DATE,  
  C CHAR(10) NOT NULL  
);

## ***More On Specifying A Relation Schema***

## ***Altering The Definition To Add Constraints***

- ALTER TABLE Plant ADD CONSTRAINT C\_20 PRIMARY KEY (P);
- ALTER TABLE Customer ADD CONSTRAINT C\_60 PRIMARY KEY (C);
- ALTER TABLE Invoice ADD CONSTRAINT C\_80 PRIMARY KEY (I);
- ALTER TABLE Customer ADD CONSTRAINT C\_70 FOREIGN KEY (P) REFERENCES Plant(P) ON DELETE SET NULL;
- ALTER TABLE Invoice ADD CONSTRAINT C\_90 FOREIGN KEY (C) REFERENCES Customer(C) ON DELETE CASCADE;



## ***Altering The Definition To Add Constraints***

- ALTER TABLE Plant ADD CONSTRAINT C\_30 UNIQUE (Pcity, Profit);
- ALTER TABLE Plant ADD CONSTRAINT C\_40 CHECK ( Pcity <> Pname );
- ALTER TABLE Plant ADD CONSTRAINT C\_50 CHECK ( (Pcity <> 'Chicago') OR (Profit > 1000) );

# ***Adding Constraint Not Currently Satisfied***

- What happens if a table is ALTERed with a CONSTRAINT that is currently not satisfied by the table?

```
drop table TEST_UNKNOWN;  
create table TEST_UNKNOWN (  
    X CHAR VARYING(10),  
    Y CHAR VARYING(10)  
);  
Insert into TEST_UNKNOWN values('a','b');  
select *  
from TEST_UNKNOWN;  
alter table TEST_UNKNOWN add constraint  
TEST_UNKNOWN check (X=Y);  
Insert into TEST_UNKNOWN values('c','d');  
select *  
from TEST_UNKNOWN;
```

# *Adding Constraint Not Currently Satisfied*

- The constraint is rejected and the system proceeds

Table dropped.

Table created.

1 row created.

| X | Y |
|---|---|
| a | b |

```
alter table TEST_UNKNOWN add constraint  
TEST_UNKNOWN check (X=Y)
```

\*

ERROR at line 1:

ORA-02293: cannot validate (KEDEM.TEST\_UNKNOWN) -  
check constraint violated

1 row created.

| X | Y |
|---|---|
| a | b |
| c | d |

# ***FOREIGN KEY Can Reference The Table Itself (We Have Seen This Before)***

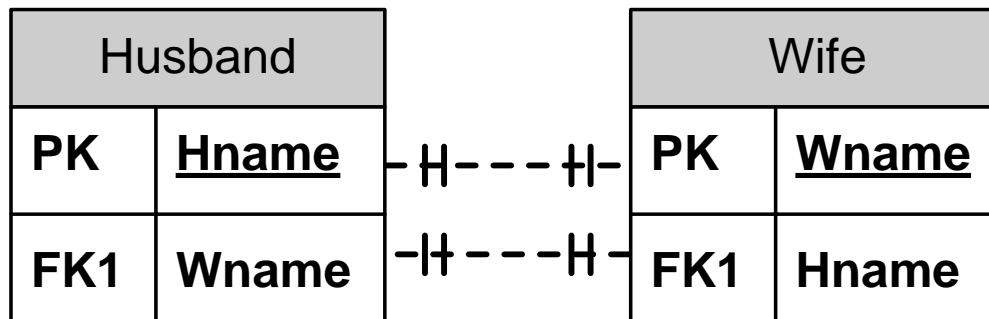
- We store information about women
- With each woman, we store her mother, if the mother is known
- The mother is a woman
- ```
create table WOMAN (  
    SSN CHAR(9) not null,  
    NAME CHAR VARYING(10) default (null),  
    MOTHER CHAR(9) default (null)  
);
```
- ```
alter table WOMAN add constraint C_01 primary key  
(SSN);  
alter table WOMAN add constraint C_02 foreign key  
(MOTHER) references WOMAN(SSN);
```

## ***Referencing Unique***

- Foreign key can also refer to UNIQUE and not only to PRIMARY KEY
- So we could also add to our database such a constraint, for which we look at an example
- ```
CREATE TABLE Test (  
  TestID CHAR(10) NOT NULL,  
  TestPname CHAR VARYING(10),  
  TestPcity CHAR VARYING(10),  
  TestProfit NUMBER  
);
```
- ```
ALTER TABLE Test ADD CONSTRAINT C_99 FOREIGN  
KEY (TestPcity, TestProfit) REFERENCES Plant(Pcity,  
Profit);
```

# ***Sometimes It Is Necessary To Define Tables First And Then Add Constraints***

- If you define a foreign key constraint, it cannot refer to a table that has not yet been designed
- Consider the following Visio diagram for some types of marriages



- You have “circular” dependencies
  - You cannot fully define Husband before Wife
  - You cannot fully define Wife before Husband
- Therefore
  1. Produce basic definitions for Husband and Wife
  2. Alter them by adding constraints later

# ***When Are Constraints Checked?***

- Essentially, each row of the TABLE has to satisfy the constraint
- Constraints are checked as tables are modified (immediately or deferred until later, generally until the end of a transaction)
- The actual checking is done either before/after each statement or at the end of a transaction
  - It is done at the end of a transaction to allow changes that cannot be done in a single statement
  - For example if  $\text{Total} = \text{Checking} + \text{Savings}$  and money is moved from Checking to Savings this constraint could be violated in the middle of the move, but must be satisfied before and after the move
- So as part of specification of a constraint one can specify
  - ***NOT DEFERRABLE*** (this is the default), or
  - ***DEFERRABLE***

# ***Assertions***



# Assertions

- Assertion is like a CHECK constraint, **but** it is not attached to a TABLE definition; it is “free floating”
- **CREATE ASSERTION** Assertion01  
CHECK  
( (SELECT COUNT (\*) FROM Plant)  
+ (SELECT COUNT (\*) FROM Customer)  
    < 1000 );
- Assertions are more natural than previously described constraints, especially when referring to several tables
- However, they are frequently not implemented
- It is very difficult to implement them both correctly and efficiently

# ***UNIQUE And PRIMARY KEY***

# ***UNIQUE And PRIMARY KEY***

- Back to our old example
- `CREATE TABLE City`  
Country NOT NULL,  
State,  
Name NOT NULL,  
Longitude NOT NULL,  
Latitude NOT NULL  
);
- A city can be identified in one of two ways
  - By its geographic location: Longitude and Latitude
  - By its official “hierarchy of names”: Country, State, Name
- It may be the case that some countries are **not** partitioned into states (or equivalent units)
  - For them it is natural to allow State to be NULL, as opposed to faking something

## ***UNIQUE And PRIMARY KEY***

- The following is OK
- `CREATE TABLE City  
Country NOT NULL,  
State,  
Name NOT NULL,  
Longitude NOT NULL,  
Latitude NOT NULL,  
UNIQUE (Country, State, Name),  
PRIMARY KEY (Longitude, Latitude) );`

# ***UNIQUE And PRIMARY KEY***

- The following ***is not OK***
- CREATE TABLE City  
Country NOT NULL,  
State,  
Name NOT NULL,  
Longitude NOT NULL,  
Latitude NOT NULL,  
PRIMARY KEY (Country, State, Name),  
UNIQUE (Longitude, Latitude) );
- ***Because State could be NULL, not permitted in primary key***
- We will see why primary keys should not contain NULLs (there are other reasons for this too)

# ***UNIQUE And PRIMARY KEY***

## ■ Small database

- CREATE TABLE City\_Population  
Country NOT NULL,  
State,  
Name NOT NULL,  
Longitude NOT NULL,  
Latitude NOT NULL,  
Population,  
PRIMARY KEY (Country, State, Name),  
UNIQUE (Longitude, Latitude) );
- CREATE TABLE City\_Size  
Country NOT NULL,  
State,  
Name NOT NULL,  
Longitude NOT NULL,  
Latitude NOT NULL,  
Size,  
PRIMARY KEY (Country, State, Name),  
UNIQUE (Longitude, Latitude) );

# ***UNIQUE And PRIMARY KEY***

- We want to combine information about cities from both tables
- ```
SELECT *  
FROM City_Population, City_Size  
WHERE (City_Population.Country = City_Size.Country  
AND City_Population.State = City_Size.State  
AND City_Population.Name = City_Size.Name) ;
```
- We will not get anything for cities in countries that are not partitioned into states!
- Because the result of comparison of say (Monaco, NULL, Monaco-Ville) = (Monaco, NULL, Monaco-Ville) is UNKNOWN
- Therefore, we cannot have (Country,State,Name) as PRIMARY KEY

# ***Workaround***

- The following can be done if we want to use UNIQUE set of attributes for joining in our example
- ```
SELECT *  
FROM City_Population, City_Size  
WHERE City_Population.Country = City_Size.Country  
AND City_Population.Name = City_Size.Name  
AND ( City_Population.State = City_Size.State  
      OR (City_Population.State IS NULL  
          AND City_Size.State IS NULL ) );
```
- But this is very burdensome and potentially easily forgotten



# ***AUTOINCREMENT for the Primary Key***

# ***Auto-incrementing Primary Keys***

- **Primary keys are critical to the maintenance of the database**
  - The primary key, in effect, identifies a specific row in a table (assume no duplicate rows, for simplicity)
  - It is used typically for binary many-to-one mappings
- Frequently, it is desirable that primary keys **do not** reflect anything in the real world
  - Do not use, Social Security Numbers, Passport Numbers, or similar
- Frequently, it is desirable that primary keys occupy little space (this is important) and that they be “meaningless”
  - To have less space devoted to foreign keys (which are typically copies of primary keys)
  - To have less space devoted to indexes as we will see in a future unit; **this is important for performance**
  - Skip to page 65.

# Oracle Version

- The simplest case is to start with 1 and increment by 1

- **Replace**

- CREATE TABLE Plant (  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
PRIMARY KEY (P)  
);

by

- CREATE TABLE Plant (  
Plant\_Id NUMBER GENERATED ALWAYS AS IDENTITY,  
P CHAR(10) NOT NULL,  
Pname CHAR VARYING(10),  
Pcity CHAR VARYING(10),  
Profit NUMBER,  
PRIMARY KEY (Plant\_Id),  
UNIQUE (P),  
);

# Oracle Version

- Plant\_Id will be **auto-generated**
- The system needs to be told that it will serve as a primary key, otherwise it will not know that
- Auto-generated primary key must be a number and a single column
- When we insert the first tuple into Plant, specifying some values of P, Pname, Pcity, Profit, then 1 will be automatically inserted into Plant\_Id
- So if we insert (276, Main, Chicago, 0.5), the system will insert (1, 276, Main, Chicago, 0.5)
- When we insert the second tuple into Plant, specifying some values of P, Pname, Pcity, Profit, then 2 will be automatically inserted into Plant\_Id
- So if we insert (93, Small, Boston, 0.6), the system will insert (2, 93, Small, Boston, 0.6)

# ***Oracle Version***

- Standard users cannot modify such auto-generated primary keys
- There are various option for the autoincrement keys
  - User can specify, e.g., start with 7 and increment by 12
  - User can specify the value of the autoincrement key for a new tuple (not ALWAYS by the systems as in the example)
- Other attributes, and not only the primary key, can also be auto-generated

# Oracle Version

```
CREATE TABLE R(  
  A_ID NUMBER GENERATED ALWAYS AS IDENTITY,  
  PRIMARY KEY (A_ID),  
  A CHAR(30) NOT NULL,  
  UNIQUE (A)  
);
```

```
INSERT INTO R (A) VALUES ('ALICE');  
INSERT INTO R (A) VALUES ('BOB');
```

```
SELECT *  
FROM R;
```

| A_ID | A     |
|------|-------|
| 1    | ALICE |
| 2    | BOB   |

- The old primary key was A

# Oracle Version

```
UPDATE R  
SET (A_ID) = 5  
WHERE A = 'ALICE' ;
```

```
SET (A_ID) = 5  
*
```

ERROR at line 2:  
ORA-32796: cannot update a generated always identity column

```
SELECT *  
FROM R ;
```

| A_ID | A     |
|------|-------|
| 1    | ALICE |
| 2    | BOB   |

- Cannot change the values of the primary key in the table

# Oracle Version

```
CREATE TABLE S (  
  B_ID NUMBER GENERATED ALWAYS AS IDENTITY,  
  B CHAR(30) NOT NULL,  
  UNIQUE (B) ,  
  C NUMBER,  
  FOREIGN KEY (C) REFERENCES R(A_ID)  
);
```

```
INSERT INTO S (B) VALUES ('CAROL') ;  
INSERT INTO S (B) VALUES ('DON') ;
```

```
SELECT *  
FROM S;
```

| B_ID | B     | C |
|------|-------|---|
| 1    | CAROL |   |
| 2    | DON   |   |

- Note foreign key attribute “to connect to” table R



# Oracle Version

```
UPDATE S
SET (C) = (SELECT R.A_ID
           FROM R
           WHERE R.A = 'ALICE')
WHERE B = 'DON' ;
```

```
SELECT *
FROM S ;
```

| B_ID | B     | C |
|------|-------|---|
| 1    | CAROL |   |
| 2    | DON   | 1 |

- “Connecting” Don to Alice

# SQL Power Architect Definition

|                            |
|----------------------------|
| R                          |
| A: INTEGER NOT NULL [ PK ] |
| B: VARCHAR                 |

**Column Properties of A**

**Source for ETL Mapping**  
None Specified

**Logical Name**  
A

**Physical Name**  
A

☒ In Primary Key

**Type**  
INTEGER

**Precision**  
☐ 0

**Scale**  
☐ 0

**Allows Nulls**  
☐ No

**Auto Increment**  
☒ Yes

**Default Value**  
☐


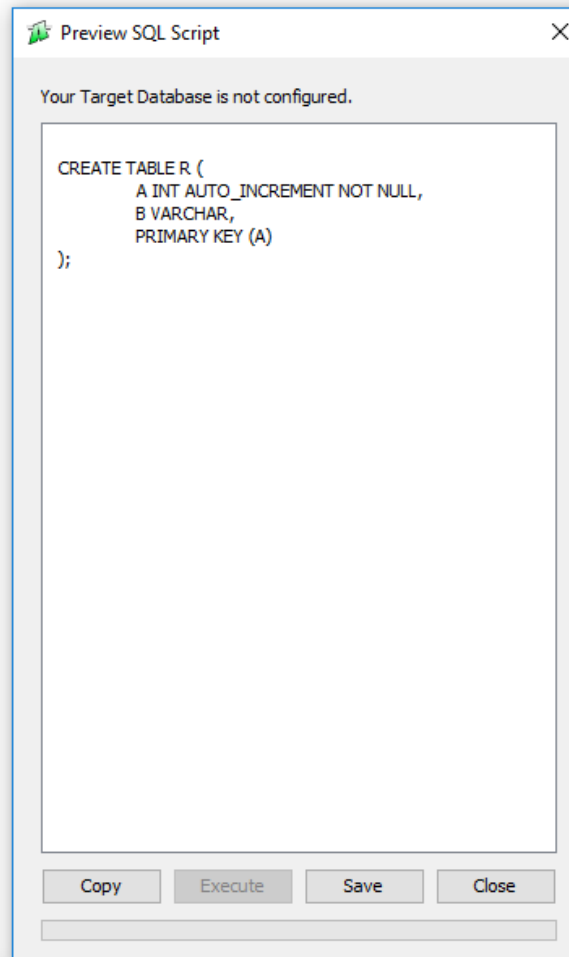
**Sequence Name (Only applies to target platforms that use sequences)**  
R\_A\_seq

**Remarks**

OK Cancel

# SQL Power Architect Generated SQL For MySQL

|                            |
|----------------------------|
| R                          |
| A: INTEGER NOT NULL [ PK ] |
| B: VARCHAR                 |



## ***More On UNKNOWN***

## ***Example On UNKNOWN Using Oracle***

- We define a table and insert tuples into it
- Our second tuple is interesting
  - For DDL and DML INSERT this satisfies the constraint (actually UNKNOWN)
  - For DML SELECT this does not satisfy the constraint (actually UNKNOWN)

```
create table TEST_UNKNOWN (  
  X CHAR VARYING(10) ,  
  Y CHAR VARYING(10) ,  
  check (X=Y)  
);
```

```
Insert into TEST_UNKNOWN values('a','a');  
Insert into TEST_UNKNOWN values('b',null);  
Insert into TEST_UNKNOWN values('c','d');
```

## *Example On UNKNOWN Using Oracle*

Table created.

1 row created.

1 row created.

```
Insert into TEST_UNKNOWN values ('c', 'd')
```

\*

ERROR at line 1:

ORA-02290: check constraint (KEDEM.SYS\_C00111308)  
violated

- Note that the third tuple does not satisfy the constraint for DDL

## ***Example On UNKNOWN Using Oracle***

```
select *  
from TEST_UNKNOWN;
```

| X | Y |
|---|---|
| a | a |
| b |   |

- The table containing tuples that satisfy the condition  $X = Y$  has two tuples

## *Example On UNKNOWN Using Oracle*

```
select *  
from TEST_UNKNOWN  
where x=y;
```

| X | Y |
|---|---|
| a | a |

- One tuple is output when condition  $X = Y$  is specified



# ***VIEWs***

# Views

- We now proceed to the definition of the user level, that is to the definition of views
- Generally speaking, a view consists of “continuously current” table that is derived by means of a **SELECT** statement from other tables
- For example, we could write

```
CREATE VIEW GoodPlant
AS SELECT *
FROM Plant
WHERE Profit > 0.0;
```
- We could now execute a query against the view

```
SELECT P
FROM GoodPlant
WHERE City = 'Chicago';
```
- This will give all P for Chicago where Profit is positive

# ***Views Versus Snapshots***

- View is not a snapshot, which is static
- View can be thought of as a macro.
- Therefore we should think of the following procedure for computing the answer to the last query:
- The system computes the value of the table GoodPlant
- The system executes the query against the table GoodPlant
- In practice, the system may compute the answer differently (e.g. as two selects), however, the result will be equivalent to the canonical procedure described above

# ***Views Defined by Queries***

- In general, almost any query definition could be used to define a view, so we could have:

```
CREATE VIEW Customer_In_The_City  
AS SELECT Cname  
FROM Plant, Customer  
WHERE Pcity = Ccity  
AND Plant.C = Customer.C;
```

- Views can also be defined **WITH CHECK OPTION**, which we will discuss later.

# *Updating Views*

- Views, in principle, can be updated just like the base tables
- However, all updates to views must be reflected in a correct update to the base table.

- Let us start with the view

```
CREATE VIEW GoodPlant  
AS SELECT *  
FROM Plant  
WHERE Profit > 0.0;
```

- Then it is clear what should be inserted into the table Plant if the following is issued:

```
INSERT INTO GoodPlant  
VALUES (675, 'Major', 'Philadelphia', .25);
```

# ***Updating Views While Forcing Defaults***

- Consider now the view

```
CREATE VIEW SomePlant  
AS SELECT P, Pname, City  
FROM Plant;
```

- Then, if the value of Profit can be NULL or has a defined default value, it is clear what should be inserted into the table Plant if the following is issued:

```
INSERT INTO SomePlant  
VALUES (675, 'Major', 'Philadelphia');
```

## ***Update To View Not Reflected In It***

- Consider the view

```
CREATE VIEW Plant_In_Chicago  
AS SELECT *  
FROM Plant  
WHERE City = 'Chicago';
```

- According to SQL the following update is valid

```
INSERT INTO Plant_In_Chicago  
VALUES (897,'Minor','Philadelphia',.1);
```

- It is reflected properly in the base table Plant, however, it does not show in the view, of course

# *Checking for Updates Not Reflected in View*

- Instead, if we define the view

```
CREATE VIEW Plant_In_Chicago  
AS SELECT *  
FROM Plant  
WHERE City = 'Chicago'  
WITH CHECK OPTION;
```

- Then the update

```
INSERT INTO Plant_In_Chicago  
VALUES (897,'Minor','Philadelphia',.1);
```

will be rejected



# ***Some Views Cannot Be Updated***

- Consider the view

```
CREATE VIEW Profit_On_Date  
AS SELECT Profit, Date  
FROM Plant, Invoice, Customer  
WHERE Plant.P = Customer.P  
AND Invoice.C = Customer.C;
```

- There is no meaning to the update

```
INSERT INTO Profit_On_Date  
VALUES (0.9,2009-02-01);
```

- Why?

- Because there is no well-defined way for reflecting this update in the base tables
- Several tables would need to be modified in a non-deterministic fashion

# ***Some Views That Cannot Be Updated***

- Consider the view

```
CREATE VIEW Avg_Amt  
AS SELECT AVG(Amt)  
FROM Invoice  
WHERE Idate = '2009-02-01';
```

- It is not permitted to issue:

```
INSERT INTO Avg_Amt  
VALUES (75);
```

- There is no way of changing the base tables in a well-defined way.

# ***Some Views That Cannot Be Updated***

- Consider the view

```
CREATE VIEW Cities_With_Plant  
AS SELECT Pcity  
FROM Plant;
```

- It is not permitted to issue

```
INSERT INTO Cities_With_Plant  
VALUES ('Palm Beach');
```

- P cannot have a NULL value, as it was the primary key

# ***Views That Are Updatable In Standard SQL***

- The following are the major conditions (there are others) that must be true for an updatable view
  - Is drawn from one TABLE  
No joins, unions, differences, intersections
  - If the underlying TABLE is a view, it must be updatable
  - The SELECTed columns are column references (each column at most once and without DISTINCT) and not values or aggregates
  - No GROUP BY
  - Skip to page 88.

# ***TRIGGERS***

## *Using A Trigger To Update A View*

- **CREATE TABLE** r (  
a CHAR (10) NOT NULL,  
b CHAR (10) NOT NULL,  
PRIMARY KEY (a)  
);
- **CREATE TABLE** s (  
a CHAR (10) NOT NULL,  
c CHAR (10) NOT NULL,  
PRIMARY KEY (a)  
);
- **CREATE VIEW** t AS  
SELECT r.a AS a, r.b AS b, s.c AS c  
FROM r, s  
WHERE r.a = s.a;

## *Using A Trigger To Update A View*

- **CREATE TRIGGER** trigger02  
**INSTEAD OF UPDATE ON** t  
**REFERENCING NEW AS** new  
**BEGIN UPDATE** s  
**SET** c = :new.c  
**WHERE** a = :old.a;  
**END** trigger02;  
  
.  
RUN
- **UPDATE** t  
**SET** c = 'q'  
**WHERE** a = '2';

# Using A Trigger To Update A View

- Tables R, S, and view T before update on the view

| A | B |
|---|---|
| 1 | e |
| 2 | f |

| A | C |
|---|---|
| 1 | m |
| 2 | n |
| 3 | o |

| A | B | C |
|---|---|---|
| 1 | e | m |
| 2 | f | n |



# Using A Trigger To Update A View

- Tables R, S, and view T after update on the view using trigger02

| A | B |
|---|---|
| 1 | e |
| 2 | f |

| A | C |
|---|---|
| 1 | m |
| 2 | q |
| 3 | o |

| A | B | C |
|---|---|---|
| 1 | e | m |
| 2 | f | q |

## ***Using A Trigger To Update (?) A View***

- Triggers will allow you to do very strange things

## *Using A Trigger To Update (?) A View*

- **CREATE TRIGGER** trigger03  
**INSTEAD OF UPDATE ON** t  
**REFERENCING NEW AS** new  
**BEGIN UPDATE** r  
**SET** b = :new.c  
**WHERE** a = :old.a;  
**END** trigger03;  
  
.  
RUN
- **UPDATE** t  
**SET** c = 'q'  
**WHERE** a = '2';

# *Using A Trigger To Update (?) A View*

- Tables R, S, and view T before update on the view

| A | B |
|---|---|
| 1 | e |
| 2 | f |

| A | C |
|---|---|
| 1 | m |
| 2 | n |
| 3 | o |

| A | B | C |
|---|---|---|
| 1 | e | m |
| 2 | f | n |

# Using A Trigger To Update (?) A View

- Tables R, S, and view T after update on the view using trigger03

| A | B |
|---|---|
| 1 | e |
| 2 | q |

| A | C |
|---|---|
| 1 | m |
| 2 | n |
| 3 | o |

| A | B | C |
|---|---|---|
| 1 | e | m |
| 2 | q | n |

# ***Modifying Objects***

# ***ALTER , DROP, REPLACE***

- In general, if an object is CREATED, it can subsequently be
  - **ALTER**ed (some features are changed)
  - **DROP**ped (removed)
- Allowing such modifications is why it is generally a good idea to name constraints, assertions, triggers, etc, while creating them

# ***Privileges***



# *Privileges*

- Privileges can be granted to user or PUBLIC for
  - Operations
  - References
  - on
  - Base tables
  - Views
- These are technically part of ***Data Control Language*** or ***DCL***

# ***Types of Privileges***

- Data-Specific
  - Select
  - Insert
  - Update
  - Delete
  - References
- Object-Specific
  - Create
  - Drop
  - Alter
- Execution
  - Execute (a PL/SQL procedure)

# ***Examples of Privileges***

- A typical instruction is:
  - **GRANT** SELECT, INSERT  
ON Customer  
TO Li, Brown;
- Privileges can be restricted to columns:
  - **GRANT** SELECT  
ON Customer.City  
TO Li, Brown;
- It is possible to grant all privileges by:
  - **GRANT** ALL  
ON Customer  
TO Li, Brown;

# ***Passing Privileges***

- It is possible to allow the users to pass the privileges to other users by issuing:
  - **GRANT** SELECT, INSERT  
ON Customer  
TO Li, Brown  
**WITH GRANT OPTION;**
- Then Li can issue
  - **GRANT** SELECT  
ON Customer.City  
TO JONES;

# ***Privilege To Reference***

- It is possible to allow a user to use columns in a table as foreign keys referring to primary keys or UNIQUE in a table to which the user has no privileges:
  - **GRANT ALL**  
ON Invoice  
TO Li;
  - **GRANT REFERENCES (C)**  
ON Customer  
TO Li;
- This privilege must be explicitly granted because Li may be able to check if a particular C appears in Customer
  - To check if C = 123456789 appears in Customer, Li attempts to INSERT an Invoice from C = 123456789
  - If C = 123456789 does not appear in Customer, the database will complain about violation of FOREIGN KEY constraint
  - If C = 123456789 appears in Customer, the database will not complain about violation of FOREIGN KEY constraint
  - This is how Li can check this and that's why this needs to be explicitly permitted if Li can insert tuples into Invoice

# ***Privilege To Reference***

- We assumed that C is a semantically meaningful number, such as SSN
  - That's why we wanted to keep it confidential whether the person with that SSN was in Customer
- If we use an autoincrement primary key the knowledge of the primary key does not provide too much useful semantic information (still: chronology + number of rows)
  - For example, the primary key could be just a row number, as we have done previously
  - Then, by adding tuples to Invoice Li can only learn (useless information) whether there is information about a customer in a particular row of Customer, without learning anything about the customer in that row
- So little information will leak even if we somebody has
  - All the privileges on Invoice
  - The privilege to reference on Customer

# ***Privileges On Views***

- It is possible to grant privileges on views.
  - Of course, the privilege must be meaningful. That is a privilege to update can be given only on a view that can be updated, etc.

# ***Revoking Privileges***

- Privileges can be revoked
- There are various way to specify what happens with privileges granted by somebody from whom a privilege is taken away



# ***Key Ideas***

# ***Key Ideas***

- CREATE for defining tables
  - Specifying domains
  - PRIMARY KEY
  - UNIQUE
  - FOREIGN KEY
  - NOT NULL
  - CHECK
  - DEFAULT
- UNKNOWNs
- Maintenance of referential integrity
- Constraint checking
  - NOT DEFERRABLE
  - DEFERRABLE
- ASSERTIONs

# ***Key Ideas***

- Auto-generated primary keys
- Triggers on INSERT, UPDATE, DELETE, firing BEFORE, AFTER, INSTEAD
- Views
- Updating a view with an SQL UPDATE
- Updating a view with an INSTEAD TRIGGER
- ALTER, DROP, REPLACE
- Privileges